

PRMS AI Processing with Amazon Bedrock

What happens to the information submitted to PRMS when we use AI, where it is processed, and who can see it

Audience: PRMS stakeholders, CGIAR Centers and Programs, management, Legal and Data Protection, and functional users

Status: Final version for stakeholder communication

Version: 3.0 (Final) 19 August 2026

Executive summary

We use Amazon Bedrock, a service from Amazon Web Services (AWS), to access the AI models used by PRMS. When we use AI, the information needed for the task is processed through AWS-managed services. We do not send this information directly to the model providers' public services or APIs.

During AI processing, the information remains within AWS-managed infrastructure. AWS states that, under the standard Amazon Bedrock configuration, model providers do not have access to the information sent to the model or the responses it generates. AWS also states that this information is not used to train the AI models.

Our AI processing is currently limited to supported AWS Regions within the United States. AWS documents some exceptions for specific models or configurations. For this reason, we review these conditions before using a new model or changing our AI configuration.

1. Why we prepared this document

PRMS allows CGIAR Centers and Programs to upload evidence of their results in formats such as PowerPoint, PDF, Excel, and Word. We use AI to read this material, find relevant information, map it to the CGIAR results framework, and extract the information we need. These files may contain unpublished results, partner information, personal data, or other sensitive information.

For this reason, there are three questions we need to answer clearly: what happens to this information, where is it processed, and can the model providers see it or use it?

We have reviewed our current AI implementation and AWS configuration to answer these questions. In this document, we explain the results of that review in simple language.

2. How we use Amazon Bedrock

2.1 Where the information goes

PRMS runs in an AWS environment managed by the Alliance infrastructure team. For our AI features, we use Amazon Bedrock, a service managed by AWS. The information needed for these tasks is processed

within AWS-managed infrastructure and is not sent directly to the model provider's public service or API. [S2]

AWS is responsible for protecting the infrastructure that supports its services. We remain responsible for how we manage our data, access, permissions, and retention periods. [S2, S10]

2.2 Step by step

Step	What happens, and where
1. Upload	The user uploads a document to PRMS.
2. Preparation	We store and prepare the file in our Alliance-managed AWS environment. Here, we read the document and extract its content.
3. AI request	We send to Amazon Bedrock only the content needed for the AI task. The original uploaded file is not sent to Amazon Bedrock.
4. AI processing	Amazon Bedrock processes the content using the selected AI model on AWS-managed infrastructure and returns a response to us.
5. Result	We review the response and use it in PRMS, for example, to show information to the user or map it to the CGIAR results framework.

The important distinction is that the uploaded file remains in our environment, and we send to Amazon Bedrock only the content needed for the AI task.

2.3 How we connect

We connect to Amazon Bedrock using AWS access and authentication controls. We do not use accounts, credentials, or API keys provided by the model providers. Permissions are managed by the Alliance within AWS and allow us to control which AI models PRMS can access. [S2]

3. Who can see and use the information

AWS states that, under the standard Amazon Bedrock configuration, model providers do not have access to the information sent to the model or the responses it generates. AWS also states that this information is not used to train the AI models. [S1, S2, S8]

Our current configuration does not allow this information to be shared with the model provider. We review these conditions before using a new model or changing our configuration. [S3, S4]

4. Where it is processed, and what is kept

4.1 Where it is processed

Our AI processing is currently limited to AWS Regions within the United States. Amazon Bedrock may process a request in different AWS Regions within the country, but the information remains within AWS

infrastructure and does not travel over the public internet between these Regions. AWS also states that the information is encrypted while it is transferred. [S5, S6]

4.2 What is kept

AWS states that, under the standard Amazon Bedrock behavior, the information sent to the model and the responses generated are not stored. In our current configuration, the optional Bedrock feature that can log this content is also disabled. [S3, S4, S7]

Separately, we keep the information PRMS needs to operate, such as uploaded files, extracted information, application logs, and backups. This information remains in our AWS environment and is managed according to Alliance retention and deletion policies.

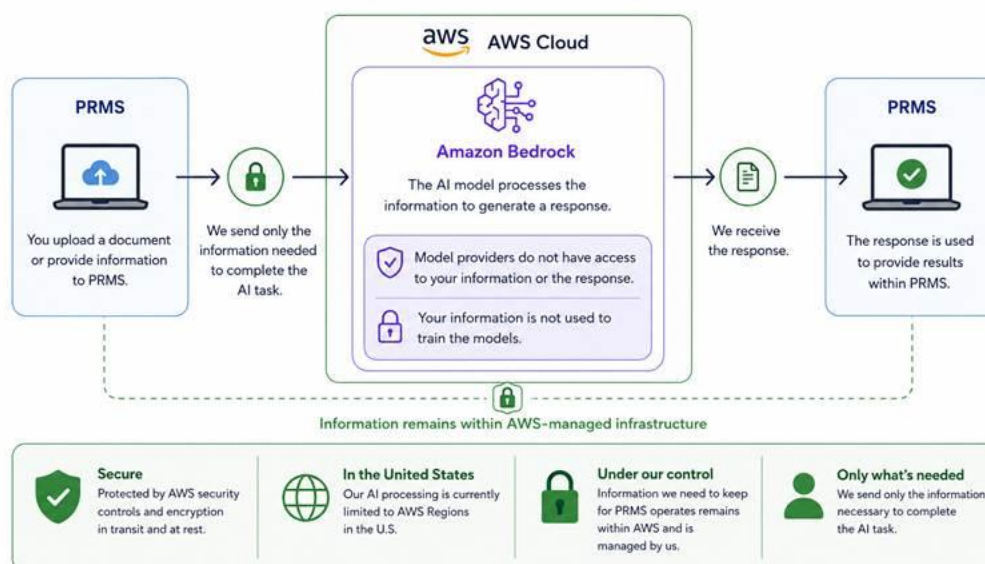
5. Our commitments

The sections above explain how Amazon Bedrock works and what AWS states about the service. In addition, we have defined our own controls to protect information when we use AI:

- We do not allow confidential or restricted information to be shared with model providers without prior review and approval from Security, Legal, and Data Protection.
- We send to Amazon Bedrock only the information needed for the AI task.
- We control access to AI models through AWS permissions and use AWS security and encryption controls to protect the information. [S9]
- We apply retention and deletion controls to the information we store. We do not keep complete documents, AI requests, or AI responses in application logs unless there is a documented need.
- We review AWS conditions before using a new model or model version, including where the information may be processed, and we update the information provided to users when needed.

How PRMS uses Amazon Bedrock

Your information stays within AWS-managed infrastructure.



6. What we tell our users

The following notice is recommended for PRMS platform pages, with a link to this document for users who want more information.

AI-assisted document processing. We use Amazon Bedrock, a service from Amazon Web Services (AWS), for our AI features. We send to Amazon Bedrock only the information needed for the AI task. During this process, the information remains within AWS-managed infrastructure.

AWS states that, under the standard Amazon Bedrock configuration, model providers do not have access to the information sent to the model or the responses it generates, and that this information is not used to train the AI models. Our AI processing is currently limited to AWS Regions within the United States.

The information we need to keep for PRMS to operate remains within AWS and is managed by us.

7. Limitations and AWS references

This document is written for a non-technical audience and is based on our review of the PRMS AI implementation and configuration, together with official AWS documentation.

The information in this document reflects our current AI configuration and the AWS conditions reviewed for the models and services we use.

This document is not legal advice, a formal data protection impact assessment, or an independent security audit.

The following official AWS sources support the information presented in this document:

- **[S1] Amazon Bedrock security, privacy, and responsible AI.** AWS states that inputs and outputs are not shared with model providers and are not used to train the base models.
<https://aws.amazon.com/bedrock/security-privacy-responsible-ai/>
- **[S2] Data protection — Amazon Bedrock.** AWS describes how model software is deployed inside the Bedrock service, the provider access boundary, and the shared-responsibility model.
<https://docs.aws.amazon.com/bedrock/latest/userguide/data-protection.html>
- **[S3] Amazon Bedrock abuse detection.** AWS describes the default behaviour of not storing or exposing inputs and outputs, and its documented exceptions.
<https://docs.aws.amazon.com/bedrock/latest/userguide/abuse-detection.html>
- **[S4] Data retention — Amazon Bedrock.** AWS documents the retention options, including the opt-in setting that allows sharing with a model provider.
<https://docs.aws.amazon.com/bedrock/latest/userguide/data-retention.html>
- **[S5] Route model inference requests across AWS Regions with cross-Region inference — Amazon Bedrock.** AWS documents geographic routing, the AWS-network path, and encryption in transit.
<https://docs.aws.amazon.com/bedrock/latest/userguide/cross-region-inference.html>
- **[S6] Regional availability by models — Amazon Bedrock.** AWS distinguishes in-Region, geographic cross-Region, and global cross-Region processing.
<https://docs.aws.amazon.com/bedrock/latest/userguide/models-region-compatibility.html>

- **[S7] Monitor model invocation — Amazon Bedrock.** AWS documents model invocation logging as an optional feature, and its storage destinations.
<https://docs.aws.amazon.com/bedrock/latest/userguide/model-invocation-logging.html>
- **[S8] Amazon Bedrock FAQs — Security.** AWS states that inputs and outputs are not shared with model providers and are not used to train AWS or third-party models.
<https://aws.amazon.com/bedrock/faqs/>
- **[S9] Data encryption — Amazon Bedrock.** AWS documents encryption in transit and key management options for supported resources.
<https://docs.aws.amazon.com/bedrock/latest/userguide/data-encryption.html>
- **[S10] AWS Data Privacy FAQs.** AWS describes customer control, shared responsibility, Regional storage choices, and the exceptions needed to provide a requested service or to comply with law.
<https://aws.amazon.com/compliance/data-privacy-faq/>